

SkyLab : An extensible workflow web application for HPC on the cloud

Vincent Paul L. Carpio
Institute of Computer Science
College of Arts and Sciences
University of the Philippines Los
Banos
vlcarpio@up.edu.ph

Katrina Joy M. Abriol-Santos
Institute of Computer Science
College of Arts and Sciences
University of the Philippines Los
Banos
khmagno@up.edu.ph

Joseph Anthony C. Hermocilla
Institute of Computer Science
College of Arts and Sciences
University of the Philippines Los
Banos
jchermocilla@up.edu.ph

ABSTRACT

Most scientific applications require high performance computing (HPC) which utilizes parallel processing to run tasks quickly and efficiently. MPI clusters are often used to cater this type of tasks but the hardware required can be costly. Peak-Two Cloud (P2C) can host HPC applications in a cloud environment which is relatively cheaper and more convenient due to on-demand provisioning. One of the key features of P2C is vCluster, a tool that can deploy MPI clusters usable through the command line. In this paper, we present the design and implementation, as well as user-evaluation results of SkyLab, a workflow web application on top of vCluster to simplify the process of running MPI applications for users not accustomed to the command line. SkyLab currently supports applications used in bioinformatics, molecular dynamics, molecular docking, and quantum chemistry. The extensible design of SkyLab enables additional tools to be incorporated easily as modules.

CCS CONCEPTS

•**Computer systems organization** → **Embedded systems**; *Redundancy*; Robotics; •**Networks** → Network reliability;

KEYWORDS

cloud computing, high performance computing, web interface

ACM Reference format:

Vincent Paul L. Carpio, Katrina Joy M. Abriol-Santos, and Joseph Anthony C. Hermocilla. 2017. SkyLab : An extensible workflow web application for HPC on the cloud. In *Proceedings of Philippine Computing Science Congress, Cebu City, Philippines, March 2017 (PCSC2017)*, 6 pages. DOI:

1 INTRODUCTION

Cloud computing has marked significant developments and possibilities in the industry. It focuses on offering services for the different needs of the modern society.

There are three categories of cloud computing services namely, Software as a Service (SaaS), Platform as a Service (PaaS) and Infrastructure as a Service (IaaS). Organizations provide SaaS depending on the demand. Google Apps is one example of SaaS that can be used to manage email and create documents, etc. PaaS offers developers a platform where they can build and deploy applications.

IaaS provides storage, servers and clusters. These tools are primarily created to serve computational needs [1]. A cloud computing platform dynamically allocates, configures, reconfigures and deallocates servers as requested or on demand. This approach ensures the elasticity of cloud computing [3].

Most scientific applications require high performance computing (HPC) which needs CPU intensive computations and large data storage. To be able to host such applications, several computers interconnected in a network such as clusters are needed. This makes scientific computing very costly in terms of hardware infrastructure investment. With the advancement in cloud computing, these scientific applications can be deployed in the cloud without worrying about hardware costs and maintenance [1]. However, studies have shown that network transmission delay is the major drawback in deploying HPC applications in the cloud[3].

Peak-Two Cloud (P2C) is a private cloud based on OpenStack designed for research in deploying HPC applications in the cloud[7]. One of the features introduced by P2C is vCluster. vCluster is a tool that enables a user to deploy a working (Message Passing Interface) cluster on demand and to terminate it after use. P2C has been used by researchers in various fields including bioinformatics, quantum chemistry, and molecular dynamics. These researchers belong to different research groups who have little or no investment in HPC infrastructure due to limited funding but requires heavy computing resources for their research. vCluster, however, is a command line application which make it difficult for non-technical users (physicists, chemists, and biologists) to use. A more user-friendly interface is needed in order to enable scientists to focus more on their science rather than on learning and using the command line.

Presented in this paper is SkyLab¹, a workflow web application on top of vCluster that addresses the concern above. Specifically, SkyLab will

- (1) allow users to execute HPC tools via web interface;
- (2) enable developers to easily extend it to support additional HPC tools;
- (3) enable users to share their instantiated clusters; and
- (4) support displaying of results using third party tools.

The following are the tools that are currently supported by SkyLab. They are commonly used by collaborators from different research groups.

- *AutoDock* - A software used to simulate protein-ligand docking[11].

PCSC2017, Cebu City, Philippines
2017. 978-x-xxxx-xxxx-x/YY/MM.
DOI:

¹<https://github.com/vincentpaul12/SkyLab>

- *AutoDock Vina* - A software similar to AutoDock 4 but on the average, it provides faster and more accurate computations [15].
- *DOCK* - Used to predict the small molecule-target interactions [9].
- *Quantum ESPRESSO* - An integrated software suite of tools for ab-initio molecular dynamics (MD) simulations and electronic structure calculations[5].
- *GAMESS* - Used for ab initio molecular quantum chemistry [12].
- *Ray* - Uses parallel genome assemblies for parallel DNA sequencing [2].
- *Impi* - MPI implementation of some image processing routines [14].

2 DESIGN AND IMPLEMENTATION

SkyLab is implemented as a web application (using the Django Web Framework²) in order to provide users access to their HPC applications by just using a web browser. This makes it challenging to implement given that multiple tools must be supported. Also, HPC applications executed through SkyLab have their own process space, separate from the process on which SkyLab is running. This makes it difficult to keep track of the applications and may even pose security threats.

Figure 1 shows the layers on which SkyLab is built on. At the bottom layer is P2C which provides the cloud infrastructure. vCluster is for on-demand provisioning and termination of MPI clusters. p2c-tools is the command line tool for activating the required HPC tool.




Figure 1: Layers on which SkyLab is built on.

The use case on which SkyLab was designed is shown in Figure 2. Users must first create a cluster then activate the desired tool to use.

2.1 User Access

Access to SkyLab is restricted to individuals with @up.edu.ph email only using django-allauth³. Tables used with socialaccount_ prefix as shown in Figure 3 are used for authentication.

2.2 MPI cluster deployment

The system spawns a thread (MPIThread) for each active cluster which handles the connection to the assigned cluster via Secure Shell (SSH). Creation and deletion of clusters is done by using vCluster commands while tool activation is done by using p2c-tools. The thread also manages task queuing and execution.

²<https://www.djangoproject.com/>

³<http://www.intenct.nl/projects/django-allauth/>

A cluster is either classified as public or private. If it is set to public, every user in the system can use it. On the other hand, for private clusters, the cluster will only be visible to the owner. The owner has the option to share the cluster to other users via the share key generated for the said cluster.

2.3 Toolsets

The system searches for Python packages inside the assigned modules folder and install it on server start. The tool sets will then be available for use with the system. The package must have a Python module named *install.py* which contains function calls for integrating the package with the system. The package must also contain the corresponding views and executable classes for each sub-tool.

2.4 Tasks

The system creates a task object for each task input by the user. A signal will then be sent and it is then received by the corresponding MPIThread which queues the task for execution. When a task is executed, it calls the assigned executable class with the given parameters. On connection error, the task waits exponentially before retrying. If the server crashes while running task execution, the task is just restarted.

Default task execution flow via executable class:

- (1) Needed remote and local directories for execution are cleared or created.
- (2) Input files are uploaded to cluster.
- (3) List of commands given are executed.
- (4) Output files are sent back to the server.
- (5) Remote task folder is deleted.
- (6) Output files are served by the server.

3 RESULTS AND DISCUSSION

3.1 System Features

The system's interface offers different functionalities that simplifies MPI cluster and task management.

- The user is authenticated by logging in with his @up.edu.ph Google account.
- The user can create an MPI cluster with optional tool activations.

Figure 4: MPI creation form

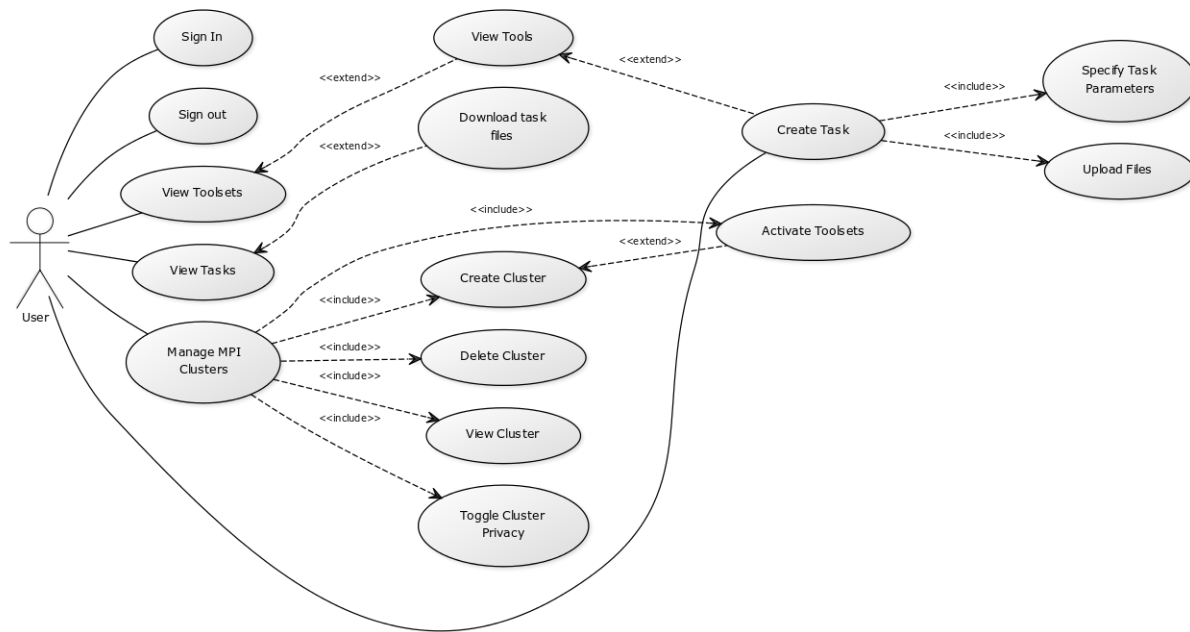


Figure 2: Use Case Diagram

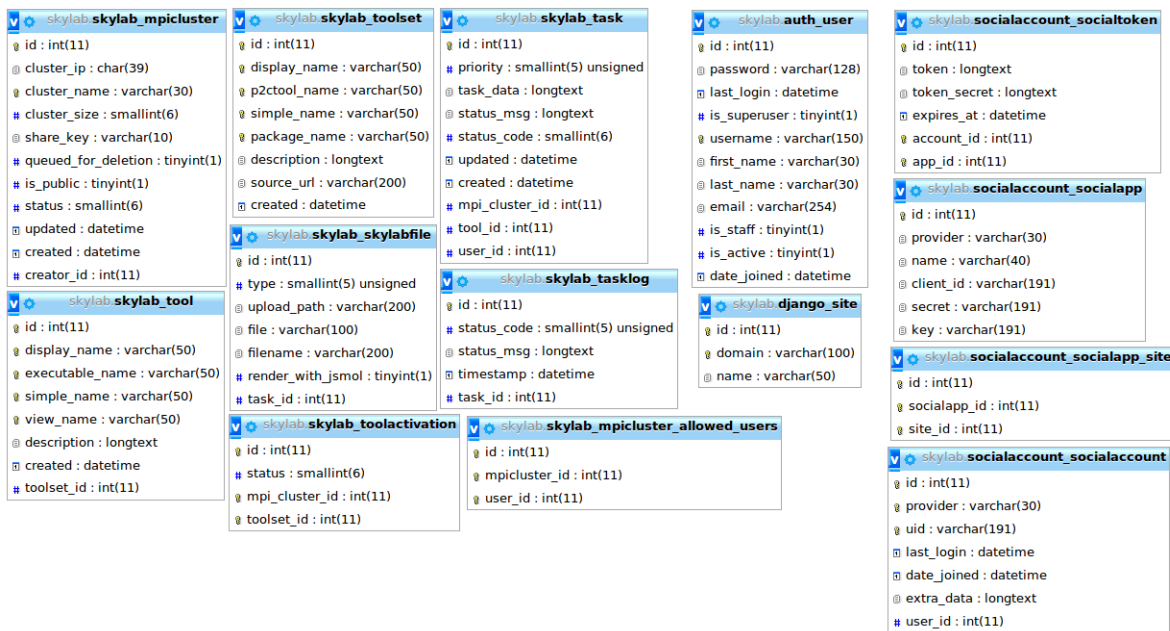


Figure 3: Database Schema

- The user can monitor visible public and private MPI clusters.

SkyLab

MCI Clusters
 Tools
 Tasks
 Joseph Anthony Hermocilla

MPI Clusters

Show entries

Search:

Cluster name	Nodes	IP address	Tasks queued	Status	Visibility	Date created
cmcsc165_image_processing	2	10.0.3.245	0	Creating	Private	01/08/17 09:27 AM

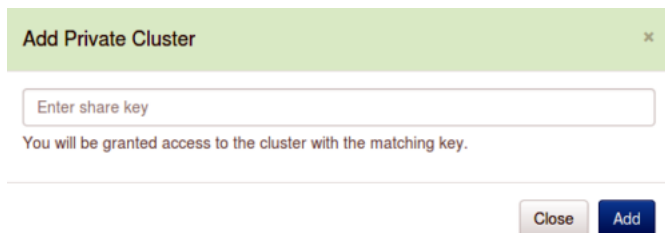
Showing 1 to 1 of 1 entries

[Previous](#)
[1](#)
[Next](#)

[Create MPI cluster](#)
[+ Add private cluster](#)

Figure 5: MPI cluster table

- The user can make a private cluster visible by entering a valid share key.



Add Private Cluster

Enter share key

You will be granted access to the cluster with the matching key.

Close Add

Figure 6: Add private cluster form

- The user can view details about a MPI cluster. If the user is the cluster's owner he has the option to delete the cluster.

cmsc165_image_processing

Cluster name	Nodes	IP address	Tasks queued	Status	Share key	Visibility	Date created
cmsc165_image_processing	2	None	0	Creating	J00NQ	Public OFF	01/08/2017 5:27 p.m.

Toolsets

Toolset	Description	Status
AutoDock 4	AutoDock is a suite of automated docking tools. It is designed to predict how small molecules, such as substrates or drug candidates, bind to a receptor of known 3D structure.	Not selected for activation Activate
Dock 6	The new features of DOCK 6 include: additional scoring options during minimization; DOCK 3.5 scoring-including Delphi electrostatics, ligand conformational entropy corrections, ligand desolvation, etc...	Not selected for activation Activate
GAMESS	The General Atomic and Molecular Electronic Structure System (GAMESS) is a general ab initio quantum chemistry package. Briefly, GAMESS can compute SCF wavefunctions ranging from RHF, ROHF, UHF, GVB, ...	Not selected for activation Activate

Figure 7: MPI detail view


- The user can select from a list which tool does he want to use.

Toolsets

AutoDock 4 AutoDock is a suite of automated docking tools. It is designed to predict how small molecules, such as substrates or drug candidates, bind to a receptor of known 3D structure.
AutoDock Vina AutoDock Vina is an open-source program for doing molecular docking. It was designed and implemented by Dr. Oleg Trott in the Molecular Graphics Lab at The Scripps Research Institute.
Dock 6 The new features of DOCK 6 include: additional scoring options during minimization; DOCK 3.5 scoring-including Delphi electrostatics, ligand conformational entropy corrections, ligand desolvation, etc...
GAMESS The General Atomic and Molecular Electronic Structure System (GAMESS) is a general ab initio quantum chemistry package. Briefly, GAMESS can compute SCF wavefunctions ranging from RHF, ROHF, UHF, GVB, ...
Impi Image processing tool that runs in parallel via MPI
Quantum ESPRESSO Quantum Espresso is an integrated suite of Open-Source computer codes for electronic-structure calculations and materials modeling at the nanoscale. It is based on density-functional theory, plane wave...
Ray Ray is a parallel software that computes de novo genome assemblies with next-generation sequencing data.

Figure 8: Toolset list view

- The user can submit a task by filling up a tool's task creation form.



Impi

MPI Cluster*
cmsc165_image_processing (nodes : 2, tasks queued: 0) (Impi status: Queued for installation)

Getting an empty list? Try creating an MPI Cluster first.

Image file(s)*
Browse... Lemna.jpg
Only supports JPEG format (.jpg, .jpeg)

Image processing operation
Sobel

+ Add operation

Figure 9: GAMESS task creation form

- The user can monitor created tasks.

Tasks

Show 10 entries Search:

Task number	Tool	MPI Cluster	Task status	Date created	Last updated
147	Impi	cmsc165_image_processing	Queued	01/08/17 09:41 AM	01/08/17 09:41 AM

Showing 1 to 1 of 1 entries

Previous 1 Next

+ Create new task

Figure 10: Task table view


- The user can view results of tasks. JSmol renders the compatible output files[6].

Task 29

GAMESS @ testcluster1 Success

Input Files:
glyc_makelp.inp

Output Files:
H2O2.log
glyc_makelp.log
task_29-scratch_files.zip


Figure 11: Task detail view

4 EVALUATION




Figure 12: Results of QUIS for SkyLab

The system has been evaluated by 56 respondents by answering a survey based on Questionnaire for User Interface Satisfaction (QUIS) [4]. Respondents are students who are unfamiliar with both HPC tools and the concept of MPI systems. Respondents are asked to test features of SkyLab by following a set of instructions and using input files provided. On the average, the users rated their overall experience with SkyLab to 6.9/10. The users listed the simplicity of the user interface to be the most positive aspect of the system while the slow speed of task processing is said to be the most negative. Majority of the tools supported by SkyLab have inherently long processing time which is not known to the respondents. The system does not focus on optimizing the said tools to achieve better performance but rather it focuses on simplifying the user's task submission process.

5 RELATED WORK

Ganglia is a system designed to monitor high performance computing systems. It uses a hierarchical model in managing the system of clusters. It uses optimized data structures and communication algorithms to achieve scalability with high concurrency. It is claimed to be used by over 500 clusters around the world. This implies that the system is tested and trusted to be used for real-world applications[10].

One of the main inspirations for developing SkyLab is the Yabi system. It provides a web interface with support for workflow environments with focus on introducing HPC applications to non-technical audience. Users can create and reuse workflows, and manage large amounts of data while system administrators can configure tools via the web interface as well. It is currently in use by multiple institutions, and is maintained as an open-source project[8].

Another related project is Web Interface for mpiBLAST (Wimpi-BLAST). It supports mpiBLAST, a parallel implementation of Basic Local Alignment Search Tool (BLAST). BLAST is a software used for sequence homology similarity search in large databases of gene sequences. mpiBLAST can utilize HPC clusters to achieve faster computing speeds but it requires knowledge in using MPI commands to benefit from its advantages. WimpiBLAST addresses this problem by providing the user a web interface to simplify the steps to use mpiBLAST[13].

6 CONCLUSION AND FUTURE WORK

The system created allows users to manage MPI clusters and submit tasks without the need for technical expertise in scripting. This makes the advantages of HPC available to non-technical users. This is achieved by parsing form inputs to generate commands for task execution. Task files can be download from the server and output files are displayed with the help of JSmol[6]. The system is also configured to install tool sets found in the modules folder making it possible to accommodate additional tools. Based on the user acceptance test conducted, the users found the system to be acceptable in terms of the criteria provided, in general.

The system achieved its main objectives but its features can still be improved and additional features can be introduced. Improved input parameter checking and error handling will make the system more robust. There are still use cases of tools that are yet to be supported. Input file generation can make the process more interactive and more customizable. Workflow design support will enable users to run complex tasks. Support for custom MPI programs will make it easier for developers to utilize the system as a test environment. Task scheduling and resource management algorithms can be used to efficiently handle resource-intensive or time consuming tasks. For example, a cluster can borrow resources from idle clusters. These recommendations will provide the users a better experience in using the system for academic and research purposes.

ACKNOWLEDGMENT

This work is supported by the Philippine Department of Science and Technology Accelerated Science and Technology Human Resource Development Program. We also would like to thank our collaborators from the Physics Department and Institute of Chemistry especially Dr. Abraham Padama and Dr. Marlon Manalo.

REFERENCES

- [1] Sanjay P. Ahuja and Sindhu Mani. 2012. The State of High Performance Computing in the Cloud. *Journal of Emerging Trends in Computing and Information Sciences* 3, 2 (Feb. 2012), 263–266. <http://www.chinacloud.cn/upload/2012-03/12031713036456.pdf>
- [2] Sbastien Boisvert, Frdric Raymond, ftlnie Godzaridis, Franois Lavolette, and Jacques Corbeil. 2012. Ray Meta: scalable de novo metagenome assembly and profiling. *Genome Biology* 13, 12 (2012), R122. DOI: <http://dx.doi.org/10.1186/gb-2012-13-12-r122>
- [3] Ivona Brandic, Ioan Raicu, Satish Narayana Srirama, Oleg Batrashev, Pelle Jakovits, and Eero Vainikko. 2011. Scalability of parallel scientific applications on the cloud. *Scientific Programming* 19, 2/3 (2011), 91 – 105. <http://search.ebscohost.com/login.aspx?direct=true&db=a9h&AN=66692024&site=ehost-live>
- [4] J. P. Chin, V. A. Diehl, and K. L. Norman. 1988. Development of an instrument measuring user satisfaction of the human-computer interface. In *Proceedings of SIGCHI '88 ACM/SIGCHI*. New York, 213–218.
- [5] Paolo Giannozzi, Stefano Baroni, Nicola Bonini, Matteo Calandra, Roberto Car, Carlo Cavazzoni, Davide Ceresoli, Guido L Chiarotti, Matteo Cococcioni, Ismaila Dabo, Andrea Dal Corso, Stefano de Gironcoli, Stefano Fabris, Guido Fratesi, Ralph Gebauer, Uwe Gerstmann, Christos Gougousis, Anton Kokalj, Michele Lazzeri, Layla Martin-Samos, Nicola Marzari, Francesco Mauri, Riccardo Mazzarello, Stefano Paolini, Alfredo Pasquarello, Lorenzo Paulatto, Carlo Sbraccia, Sandro Scandolo, Gabriele Sclauzero, Ari P Seitsonen, Alexander Smogunov, Paolo Umari, and Renata M Wentzcovitch. 2009. QUANTUM ESPRESSO: a modular and open-source software project for quantum simulations of materials. *Journal of Physics: Condensed Matter* 21, 39 (2009), 395502 (19pp). <http://www.quantum-espresso.org>
- [6] Robert M. Hanson, Jaime Prilusky, Zhou Renjian, Takanori Nakane, and Joel L. Sussman. 2013. JSmol and the Next-Generation Web-Based Representation of 3D Molecular Structure as Applied to Proteopedia. *Israel Journal of Chemistry* 53, 3-4 (2013), 207–216. DOI: <http://dx.doi.org/10.1002/ijch.201300024>

- [7] Joseph Anthony C. Hermocilla. 2014. P2C: Towards Scientific Computing on Private Clouds. In *Proceedings of the National Conference on Information Technology Education (NCITE 2014)*. 162–167.
- [8] Adam A. Hunter, Andrew B. Macgregor, Tamas O. Szabo, Crispin A. Wellington, and Matthew I. Bellgard. 2012. Yabi: An online research environment for grid, high performance and cloud computing. *Source Code for Biology and Medicine* 7, 1 (2012), 1 – 10. <http://search.ebscohost.com/login.aspx?direct=true&db=a9h&AN=74110216&site=ehost-live>
- [9] P. T. Lang, S. R. Brozell, S. Mukherjee, E. T. Pettersen, E. C. Meng, V. Thomas, R. C. Rizzo, D. A. Case, T. L. James, and I. D. Kuntz. 2009. DOCK 6: Combining Techniques to Model RNA-Small Molecule Complexes. *RNA* 15 (2009), 1219–1230.
- [10] Matthew L. Massie, Brent N. Chun, and David E. Culler. 2004. The ganglia distributed monitoring system: design, implementation, and experience. *Parallel Comput.* 30, 7 (2004), 817 – 840. <http://search.ebscohost.com/login.aspx?direct=true&db=a9h&AN=13956548&site=ehost-live>
- [11] G. M. Morris, R. Huey, W. Lindstrom, M. F. Sanner, R. K. Belew, D. S. Goodsell, and A. J. Olson. 2009. Autodock4 and AutoDockTools4: automated docking with selective receptor flexibility. *J. Computational Chemistry* 2009 (2009), 2785–91.
- [12] Michael W. Schmidt, Kim K. Baldridge, Jerry A. Boatz, Steven T. Elbert, Mark S. Gordon, Jan H. Jensen, Shiro Koseki, Nikita Matsunaga, Kiet A. Nguyen, Shujun Su, Theresa L. Windus, Michel Dupuis, and John A. Montgomery. 1993. General atomic and molecular electronic structure system. *Journal of Computational Chemistry* 14, 11 (1993), 1347–1363. DOI: <http://dx.doi.org/10.1002/jcc.540141112>
- [13] Parichit Sharma and Shrikant S. Mantri. 2014. WImpiBLAST: Web Interface for mpiBLAST to Help Biologists Perform Large-Scale Annotation Using High Performance Computing. *PLoS ONE* 9, 6 (2014), 1 – 13. <http://search.ebscohost.com/login.aspx?direct=true&db=a9h&AN=96861207&site=ehost-live>
- [14] Axel S. Trajano and Joseph Anthony C. Hermocilla. 2010. *Implementation of selected image processing routines for single and distributed processors*. Technical Report. Undergraduate Special Problem, University of the Philippines, Los Banos.
- [15] Oleg Trott and Arthur J. Olson. 2010. AutoDock Vina: Improving the speed and accuracy of docking with a new scoring function, efficient optimization, and multithreading. *Journal of Computational Chemistry* 31, 2 (2010), 455–461. DOI: <http://dx.doi.org/10.1002/jcc.21334>